

BUILDARENA: A PHYSICS-ALIGNED INTERACTIVE BENCHMARK OF LLMs FOR ENGINEERING CONSTRUCTION

Tian Xia¹, Tianrun Gao¹, Wenhao Deng¹, Long Wei¹
 Xiaowei Qian¹, Yixian Jiang¹, Chenglei Yu¹, Tailin Wu^{1*}
¹Westlake University
 {xiatian20, wutailin}@westlake.edu.cn



Figure 1: Examples of **BuildArena**'s construction results by LLMs, covering three tasks: **Lift** (left subfigure), **Transport** (upper right), and **Support** (lower right).

ABSTRACT

Engineering construction automation aims to transform natural language specifications into physically viable structures, requiring complex integrated reasoning under strict physical constraints. While modern LLMs possess broad knowledge and strong reasoning capabilities that make them promising candidates for this domain, their construction competencies remain largely unevaluated. To address this gap, we introduce **BuildArena**, the first physics-aligned interactive benchmark designed for language-driven engineering construction. It contributes to the community in four aspects: (1) a highly customizable benchmarking framework for in-depth comparison and analysis of LLMs; (2) an extendable task design strategy spanning static and dynamic mechanics across multiple difficulty tiers; (3) a 3D Spatial Geometric Computation Library for supporting construction based on language instructions; (4) a baseline LLM agentic workflow that effectively evaluates diverse model capabilities. On eight frontier LLMs, **BuildArena** comprehensively evaluates their capabilities for language-driven and physics-grounded construction automation. The project page is at build-arena.github.io.

*Corresponding author.

1 INTRODUCTION

Engineering construction automation is an important field of AI for Engineering. It has various applications in domains such as automotive, transportation, and civil infrastructure (Wang et al., 2005; Domingues et al., 2016; Lin et al., 2019). The goal is to translate high-level task descriptions into executable end-to-end build plans that cover design, fabrication, and assembly. An ideal workflow lets users describe what they want in plain terms. For example, users request “*Design a rocket that meets Mars mission requirements.*” The system then creates realistic parts with precise material details and manufacturing specs. The workflow also provides assembly instructions that can be integrated into production systems. Such automation capabilities promise significant improvements in engineering efficiency and productivity.

Language-driven automated construction presents challenges in two key aspects. On one hand, it necessitates physics simulation environments with high fidelity to real-world constraints so that virtual designs and assembly procedures adhere to geometric, physical, and structural constraints. While modern physics engines and robotics benchmarks offer robust simulation capabilities (Todorov et al., 2012; Coumans & Bai, 2016; Makoviychuk et al., 2021), there remains a gap in environments that integrate physics verification with language-driven multi-component assembly processes. On the other hand, the domain demands multilevel reasoning across long temporal horizons and 3D spatial contexts, as engineering artifacts inherently exhibit hierarchical organization, and their assembly follows sequential dependencies with strict feasibility constraints (Jiménez, 2013; De Fazio & Whitney, 2003). These factors collectively require both breadth of domain knowledge and depth of analytical thinking, challenging even for expert human engineers.

Large language models (LLMs) have progressed rapidly in recent years, accumulating broad world knowledge and demonstrating strong capabilities in language understanding, mathematical reasoning, and code generation (Brown et al., 2020; Guo et al., 2025; Shao et al., 2024; Roziere et al., 2023). Moreover, they have shown proficiency in following human instructions, generating plans, invoking tools, and composing executable programs to interact with the external environments (Yao et al., 2023; Schick et al., 2023; Qin et al., 2023). These general intelligent capabilities position LLM agentic systems as promising candidates for automatic engineering construction.

Despite these advances, current evaluations on LLMs provide insufficient evidence of their capacity to construct physical entities. Established LLM benchmarks predominantly assess mathematical and programming capabilities (Cobbe et al., 2021; Hendrycks et al., 2021; Chen et al., 2021), which are evaluated mainly in textual or static environments, without interactions with physical environments. Existing physical reasoning datasets focus on physics understanding, but neglect the multi-step construction processes (Bakhtin et al., 2019; Cherian et al., 2024). Meanwhile, research in programmatic 3D or CAD generation has advanced generation performance but rarely validates whether the generated designs yield executable assemblies under realistic physical conditions (Jones et al., 2020; Mallis et al., 2024). This interdisciplinary gap highlights the absence of frameworks to evaluate whether LLMs can effectively translate natural language specifications into physically viable assemblies. This limitation motivates our **research question**: *How can we comprehensively evaluate LLMs for language-driven and physics-grounded construction automation?*

In this paper, we answer the research question by proposing **BuildArena**, a physics-aligned interactive benchmark designed to assess LLMs’ capabilities in engineering construction tasks. To our knowledge, this is the first benchmark that enables LLMs to perform 3D structure construction via natural language instructions and evaluates their performance within a physically constrained environment. Examples of construction results are shown in Figure 1. Comparison between **BuildArena** and existing benchmarks are provided in Table 1 (*See Appendix A for more related work*), implying that our work substantially expands the scope of current LLM benchmarks to 3D construction domains. Our technical contributions are summarized as follows.

- **We establish a highly customizable benchmarking framework for in-depth comparison and analysis of LLMs, featuring detailed logging.** It by default consists of three components: task definition, LLM-based construction, and simulation-based evaluation. It supports customization of each component (see Figure 2).
- **We create an extendable task design strategy.** The strategy includes three task categories with quantifiable difficulty levels, along with corresponding evaluation metrics.

Table 1: Comparison between **BuildArena** and previous benchmarks.

Benchmarks	Spatial Reasoning	3D Construction	Construction-aimed Planning	Physical Simulator	Interactive Environment
PlanBench (Valmeekam et al., 2023)	✗	✗	✗	✗	✗
PlanQA (Rodionov et al., 2025)	✓	✗	✗	✗	✗
PHYRE (Bakhtin et al., 2019)	✓	✗	✗	✓	✓
VOYAGER (Wang et al., 2023a)	✓	✗	✓	✓	✓
Embodied Agent Interface (Li et al., 2024)	✓	✓	✗	✓	✓
BuildArena (ours)	✓	✓	✓	✓	✓

- **We develop a key framework module: a 3D Spatial Geometric Computation Library.** The module facilitates computations and feedback in iterative 3D construction to ensure accurate execution of LLMs’ language instructions. As the geometric computation functions in the widely used Besiege simulator (Spiderling, 2018) are closed-source and inaccessible, our open-source module reproduces its building operations.
- **We design an LLM agentic workflow that effectively evaluates diverse model capabilities.** It serves as a baseline for workflow development in the engineering construction domain.

2 METHOD

This section details our benchmarking methodology, which includes task setup in Section 2.1, language-driven and physics-grounded construction in Section 2.2, the LLM agentic workflow in Section 2.3, and the evaluation methodology in Section 2.4. Our method is illustrated in Figure 2.

2.1 TASK

To design tasks in a principled manner, we first abstracted a set of difficulty dimensions that are commonly encountered in engineering practice:

- **Quantification:** Extent of explicit numerical reasoning required (Kamble et al., 2024).
- **Robustness:** Tolerance to single-point failures (Zhao et al., 2023).
- **Magnitude:** Structural scale in span, load, and module count (Fan et al., 2023).
- **Compositionality:** Required depth of hierarchical substructure construction and integration (Thurairajah et al., 2023).
- **Precision:** Strictness of geometric requirement for placement and orientation (Gao et al., 2024b).
- **Ambiguity:** Clarity and completeness of task guidance (Moon et al., 2024).

Upon these dimensions, we construct three representative engineering task categories: **Support** (static structural stability), **Transport** (dynamic horizontal movement), and **Lift** (dynamic vertical movement). Within each task category, we defined three levels of difficulty, Easy (Lv.1), Medium (Lv.2), and Hard (Lv.3), by adjusting task details so that the corresponding requirements align with the above dimensions (see Figure 3). This design ensures both diversity of engineering scenarios (statics vs. dynamics, horizontal vs. vertical motion) and systematic coverage of engineering difficulty dimensions. Task descriptions, performance indicators and evaluation criteria of these three tasks are as follows. Detailed task content as LLM prompts are provided in Appendix F.

Transport focuses on constructing a mechanical structure capable of directional movement on a planar surface. It primarily examines the LLM agents’ ability to exploit the spatial movement afforded by given building components. As the difficulty increases beyond Lv.1, the explicit instruction of building a four-wheeled vehicle is removed, and the transportation target changes from the machine itself into a cargo load with level-specific scale, adding challenges on both instruction interpretation and building larger structures. The **maximum transport distance** is chosen as the

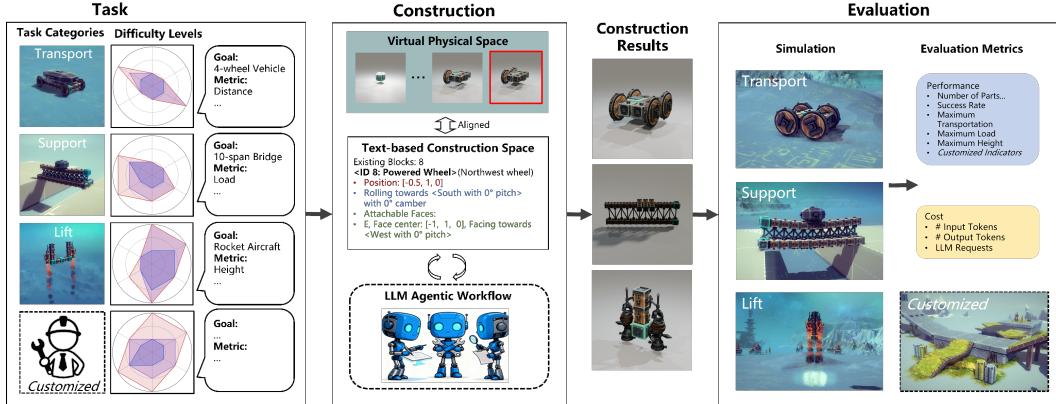


Figure 2: Illustration of our **BuildArena** framework. It contains three parts: (1) Task definition; (2) LLM-based Construction; (3) Simulation-based Evaluation. The arrows represent our pipeline. Components in dashed boxes, *i.e.*, task type, LLM agentic workflow, and simulator, could be customized by users. Details of the construction procedure is shown in Figure 4.

performance indicator, and we use a distance threshold as the criteria to identify if the machine is able to deliver effective transportation to the target.

Support requires constructing a static structure to support a load across a gap, aiming to test the ability of LLM agents to design and build bridges. The span of the gap multiplies across three levels, as larger span directly requires larger scale of the bridge, making the stable support harder as well. While only one is allowed in Lv.1, both Lv.2 and Lv.3 permit the modular construction with no more than three substructures without any detailed instruction, which also requires more precise assembly. We select **maximum load weight** as the performance indicator for bridges, and use a minimum threshold to determine if the bridge successfully supports the load.

Lift requires constructing a rocket. At Lv.1, LLMs are explicitly required to build a single rocket engine without instruction on how to build, with **thrust-to-weight ratio** (TWR) as the performance indicator. TWR > 1 represents the feasibility of providing effective thrust and marks successful construction. At Lv.2, the task requires LLMs to construct a rocket-powered aircraft as an integrated single structure. At Lv.3, LLMs must first build two separate substructures (a rocket engine and a support frame) before assembling the two into an aircraft. Both Lv.2 and Lv.3 tasks require that the aircraft are capable of launching from the ground. **Maximum Height** is adopted as the performance indicator, and the aircraft must reach a specific elevation to meet the success criterion. The escalation from Lv.1 to Lv.3 compounds multiple sources of engineering difficulty: higher demands on precise module alignment, the presence of multiple single points of failure (engine placement, structural balance), and strict requirements on modular construction and assembly. Together, these factors make **Lift** the most challenging task.

Task Customization. Tasks can be customized via textual prompts that define task constraints, objectives, testing procedures, and evaluation metrics. The prompts are fed into the LLM agentic workflow (see Section 2.3), which then executes the construction process.

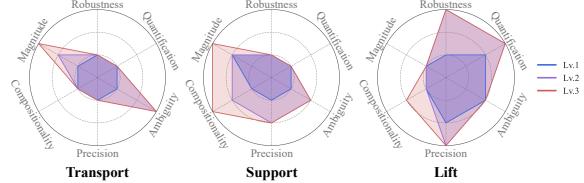


Figure 3: Difficulty profiles of the three tasks **Transport**, **Support**, and **Lift** across six engineering dimensions: Quantification, Redundancy, Scale, Modularity, Precision, and Ambiguity. Each radar chart illustrates how difficulty escalates from Lv.1 (blue) to Lv.2 (purple) and Lv.3 (red).

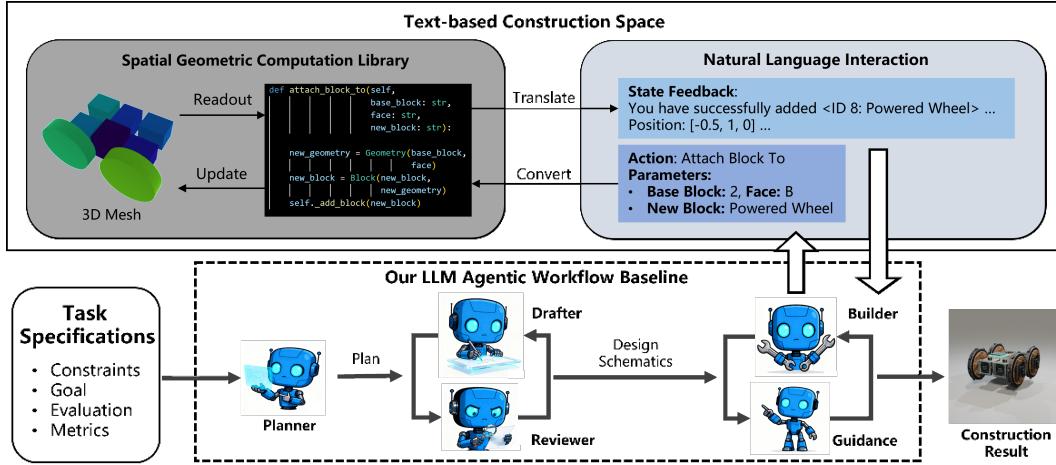


Figure 4: Details of the construction procedure in Figure 2. Our designed workflow (bottom row) contains five collaborative LLM entities, serves as a baseline for future user-customized alternatives. The text-based construction space (top row) has two transferable formats: code for physics-aligned spatial geometric computation, and natural language for LLM interface compatibility.

2.2 LANGUAGE-DRIVEN AND PHYSICS-GROUNDED CONSTRUCTION

From the perspective of human engineering practice, construction is inherently an incremental and constraint-driven process. Structures are assembled step by step, each new component must connect to existing ones, and physical feasibility (e.g., collision avoidance) is continuously verified (Wilson & Latombe, 1994; De Fazio & Whitney, 2003). Each successful action requires accurate reasoning about the spatial relationships between new and existing structures. These features necessitate *Besiege* (Spiderling, 2018), an ideal platform to evaluate the LLMs for physics-grounded construction automation. *Besiege* is a popular construction sandbox game with realistic physics simulation, widely validated by world-wide player community to align with human physical intuition. It has a rich modules space, a complete collection of basic structural and functional module types that can be combined to build complex objects, all completed by iteratively attachment and refinement of the native modules. Details about modules are provided in Appendix E.1.

However, *Besiege* only provides graphical representation of 3D structures in the interactive construction space for human users, instead of natural-language illustration for LLMs. It also only supports direct manipulation through physical controller inputs, without any interface for symbolic or language-based interaction, nor any indirect API for programmatically operating the construction process. Therefore, we develop an *open-source Spatial Geometric Computation Library* that *mimics Besiege’s closed-source construction logic and physical constraints*, enabling LLMs to interact with the construction space through language interfaces. It ensures consistency between the effects of actions executed by LLMs in the language space and those performed by human users in the graphic interface, as illustrated in Figure 4.

In implementation, it accepts an action announcing the operation and arguments from the LLM agent, computes and updates the state accordingly, and conducts physical constraint checks: it either returns a human-interpretable description of the current state, or prohibits the invalid action if certain constraints are violated and returns a prompt explaining the failure reason. All actions fall into four categories *Build*, *Refine*, *Query*, and *Control*, detailed in Appendix E.2.

2.3 LLM AGENTIC WORKFLOW

LLMs execute the construction procedure via LLM agentic workflow (Zhang et al., 2025a). For clear comparison between different LLMs, we restrict our consideration to workflows where all its entities employ the same LLM, differentiated by their respective prompts, as illustrated in Figure 2. After careful refinements, we finally obtained an effective workflow. Its design follows two principles: (1) a coarse-to-fine structure with an outline plan and progressing to granular, executable

details (Xue et al., 2024); (2) a multi-party debate and multi-turn revision framework for incremental improvement of construction quality (Du et al., 2023). Based on these principles, our workflow employs five entities: Planner (P), Drafter (D), Reviewer (R), Builder (B), and Guidance (G). In addition, Controller (C) is used for the task **Transport**. Prompt details are provided in Appendix G. The workflow includes three stages, as shown in the bottom row of Figure 4:

- **Plan Phase:** Executed by Planner, this phase takes the task description and initial module set as input, outputting a structured construction plan in a predefined format.
- **Draft–Review Loop:** Based on the generated plan, Drafter produces design schematics. Reviewer reviews and verifies the schematics, guiding Drafter’s revisions. The loop repeats until approval; and terminates in failure if the plan violates predefined rules.
- **Build–Guidance Loop:** With approved schematics as input, Builder and Guidance collaborate on execution suggestions, building actions, and feedback. Guidance generates high-level suggestions step by step based on the draft, specifying the next action to invoke; Builder converts them into formatted construction commands for the Spatial Geometric Computation Library, which updates the states and returns either descriptive feedback or error prompts. The loop ends when Guidance confirms full completion, with the final state converted via the library into a conflict-free, simulation-compatible runnable file as the final result. Rejection by Guidance based on predefined rules terminates the process in failure.

This workflow serves as a baseline for the development of advanced workflows in the future. The workflow component can be substituted with any user-customized one.

2.4 EVALUATION

Our evaluation strategy is as follows. For each task-LLM pair, run the aforementioned construction procedure to produce a result (*e.g.*, a rocket) with detailed logs (*e.g.*, token consumption, conversation turns). This result is then placed in the simulation environment to operate, yielding evaluation metrics. To enhance reliability, the procedure is sampled 64 times for each task-LLM pair, with final reported results averaged over these 64 runs.

Simulation environments are based on the Besiege sandbox, with task-specific simulation protocols. For **Transport** tasks, the system evaluates if a constructed machine achieves effective motion. The LLM agent must specify a control configuration and sequence—invalid or missing controls cause immediate failure. The machine is then loaded into the environment. If repeated attempts show no effective movement, the system concludes the structure lacks mobility. For **Support** tasks, the environment provides fixed obstacles with varying gap widths according to difficulty level. A payload of gradually increasing weight is placed on the structure, and the simulation measures whether the machine can support and stabilize the load without collapse or loss of balance. For **Lift** tasks, Lv.1 records water cannons’ heating status to calculate TWR; Lv.2–3 continuously activates water cannon firing to simulate launch, with module trajectories and cannons’ heating status recorded for evaluation during a fixed window.

Simulation environments can also be customized according to user-defined tasks. Specifically, users set test conditions for the construction result—for instance, adding a load module above a constructed bridge—and configure simulation parameters, including the initial position, tracking points, and control information. All these configurations are implemented by invoking our developed Spatial Geometric Computation Library. Finally, **BuildArena** executes the simulation using a unified script and collects log data throughout the entire simulation process.

Evaluation metrics cover performance and cost. **Performance** includes three metrics: (1) Number of parts, referring to the count of modules present in the construction result. A smaller value is preferable because a core principle in engineering prioritizes simpler system structures to improve maintainability and reliability (El-Rayes & Khalafallah, 2005). (2) Success rate, defined as the proportion of trials that successfully passed the criteria among 64 samples, a higher value is better. (3) Performance indicator, a task-specific metric extracted from simulation data that evaluates the performance under realistic physical conditions. A higher value is preferable for all indicators. Detailed success criterion and indicator setup of each task are specified in Section 2.1. **Cost** is evaluated using three metrics: (1) number of accumulated input tokens, (2) number of output tokens, and (3) total number of LLM requests. A lower value is preferable for all the cost metrics.

Table 2: Average ($n = 64$) performance comparison on different tasks across task levels Lv.1 (easy), Lv.2 (medium), and Lv.3 (hard). The indicator means maximum displacement for **Transport**; maximum load for **Support**; TWR for Lv.1 and maximum height for Lv.2, Lv.3 of **Lift**. The best results are in **bold** and the second-best are underlined.

Task	Model	Number of Parts			Success Rate (%)↑			Indicator↑		
		Lv.1	Lv.2	Lv.3	Lv.1	Lv.2	Lv.3	Lv.1	Lv.2	Lv.3
Transport	GPT-4o	11.0	13.1	18.4	9.4	<u>1.6</u>	7.8	13.5	4.2	5.1
	Claude-4	9.5	13.4	26.1	<u>17.2</u>	4.7	15.6	34.9	6.8	7.3
	Grok-4	12.0	15.6	34.6	25.0	0.0	<u>9.4</u>	<u>19.6</u>	4.2	<u>12.3</u>
	Gemini-2.0	9.0	12.3	10.0	1.6	<u>1.6</u>	1.6	4.8	<u>4.6</u>	4.6
	DeepSeek-3.1	9.2	11.6	16.7	6.2	0.0	1.6	6.2	<u>4.6</u>	4.6
	Qwen-3	9.1	7.8	10.9	10.9	<u>1.6</u>	4.7	18.4	3.3	21.5
	Kimi-K2	14.2	17.1	19.5	12.5	0.0	1.6	7.4	4.5	3.5
	Seed-1.6	7.4	11.2	<u>28.5</u>	6.2	4.7	7.8	4.3	3.8	6.8
Support	GPT-4o	36.8	16.7	29.9	40.6	0.0	0.0	181.2	0.0	0.0
	Claude-4	8.0	21.4	31.5	7.8	1.6	0.0	36.8	3.3	0.0
	Grok-4	18.7	22.3	33.3	46.9	15.6	0.0	211.4	44.5	0.0
	Gemini-2.0	20.5	23.5	<u>41.0</u>	23.4	0.0	0.0	105.5	0.0	0.0
	DeepSeek-3.1	19.0	10.5	<u>17.8</u>	25.0	0.0	0.0	122.6	0.0	0.0
	Qwen-3	18.6	18.0	22.2	12.5	4.7	0.0	70.5	13.9	0.0
	Kimi-K2	23.3	<u>34.7</u>	16.5	29.7	4.7	0.0	122.6	18.6	0.0
	Seed-1.6	33.4	36.2	68.8	<u>45.3</u>	9.4	3.1	197.4	<u>25.8</u>	7.1
Lift	GPT-4o	4.0	7.4	<u>5.1</u>	7.8	<u>3.1</u>	0.0	0.9	9.1	<u>4.1</u>
	Claude-4	4.5	<u>7.9</u>	2.5	<u>10.9</u>	1.6	0.0	<u>1.0</u>	4.1	1.2
	Grok-4	<u>4.8</u>	6.8	1.1	31.2	31.2	3.1	1.8	890.6	86.5
	Gemini-2.0	4.3	6.3	4.9	0.0	0.0	0.0	0.5	2.8	0.8
	DeepSeek-3.1	4.5	6.9	1.1	<u>10.9</u>	0.0	0.0	<u>1.0</u>	3.3	0.6
	Qwen-3	3.5	7.2	3.2	<u>3.1</u>	0.0	0.0	<u>0.6</u>	2.7	0.8
	Kimi-K2	5.3	15.1	12.0	6.2	<u>3.1</u>	0.0	0.7	<u>44.8</u>	1.8
	Seed-1.6	3.5	3.3	0.0	6.2	0.0	0.0	0.9	1.7	0.0



Figure 5: Example of the construction process. The rocket is constructed by Grok-4 for the **Lift** task under the Lv.2 (Medium) difficulty level. More examples are presented in Figure 10.

3 EXPERIMENTS

In the experiments, we aim to answer the following two questions: (1) Whether **BuildArena** serves as an effective benchmark for testing the construction capabilities of LLMs? (2) How existing mainstream models perform within the **BuildArena** framework? To answer these questions, we evaluate eight closed-source LLMs in **BuildArena**, including GPT-4o, Claude-4, Grok-4, Gemini-2.0, DeepSeek-3.1, Qwen-3, Kimi-K2, and Seed-1.6. All simulations are conducted on Besiege. Model snapshots, module space, and simulation details are provided in Appendix D. We provide the code of **BuildArena** in this link.

3.1 EFFECTIVENESS OF **BUILDARENA**

The performance of eight models on **BuildArena** is presented in Table 2, with examples of construction results shown in Figure 9 and examples of construction procedures illustrated in Figure 5. These results demonstrate that, supported by the **BuildArena** evaluation framework, LLMs achieve language-based 3D construction automation, as evidenced by the following aspects. (1) Regarding task design, the diversity and difficulty levels are reasonably configured. Across individual

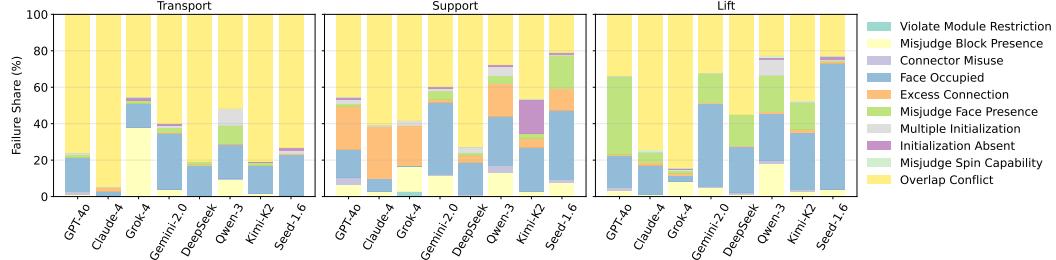


Figure 6: Distributions of failure reasons averaged over different LLMs.

tasks, performance tends to decrease as difficulty increases. An exception is the **Lift** task, where Lv.1 uses different metrics from Lv.2/3, making direct comparisons inappropriate. Specifically, at the Hard difficulty level of three tasks, most models exhibit low performance, yet a small number outperform others, indicating that the difficulty and criteria settings possess good discriminative power. (2) Concerning the LLM agentic workflow, numerous successful construction outcomes validate its effectiveness. This workflow enables collaborative behaviors among LLMs such as reflection (*e.g.*, the third subfigure from the left in Figure 5), which is essential for long-sequence planning. (3) Our Spatial Geometric Computation Library facilitates language-driven manipulation of the physical world. As illustrated in the construction procedure figures, these processes involve diverse actions including attachment, removal, rotation, shifting, and connection, which collectively meet the action requirements of construction tasks. (4) The simulator provides environmental support for the evaluation phase. For instance, it can place loads on bridges to test their load-bearing capacity and offer conditional support for the entire launch process of rockets. Overall, these key components of **BuildArena**—task design, workflow, library, and simulator—collectively provide robust support for evaluation, enabling it to function as an effective and reliable benchmark.

3.2 PERFORMANCE OF LLMs

3.2.1 LIMITED PERFORMANCE OF CURRENT LLMs

From a complementary perspective, leveraging our **BuildArena** framework, current LLMs demonstrate elementary construction capabilities, as shown in Table 2. These capabilities are reflected in three aspects as follows. (1) in **Transport** tasks, as difficulty increases from Lv.2 to Lv.3—corresponding to an increase in payload size—all models adapt by scaling up the number of components to meet size constraints, thereby maintaining moving stability during simulation. Such patterns indicate that current LLMs effectively address challenges related to magnitude and ambiguity. (2) Second, when explicit constraints are relaxed, LLMs attempt unconventional solutions: they propose propulsion-powered carriers for **Transport** tasks and wheel-integrated bridge structures for **Support** tasks. In the latter case, LLMs explicitly state the need to utilize the automatic braking function of wheels to stabilize the bridge (see Figure 10, **Support** Lv.1 by DeepSeek-3.1). These behaviors highlight the potential of LLMs for creative exploration. (3) Remarkably, structures mirroring real-world engineering practices are constructed by LLMs, such as steel trusses in bridges and differential steering in vehicles (see Figure 10, **Support** Lv.2 by Grok-4). This suggests that structural concepts learned from text are not purely symbolic but carry implicit spatial information, enabling LLMs to instantiate them as feasible 3D structures. Notably, LLMs construct structures that align with real-world engineering practices, such as widely used steel trusses in bridges and differential steering systems in vehicles. This observation suggests that structural concepts learned from text are not purely symbolic but carry implicit spatial information, enabling LLMs to instantiate them as feasible 3D structures.

However, these models still suffer from significant limitations. (1) In hierarchical assembly tasks, such as the **Support** task, LLMs’ success rates drop sharply as the assembly complexity, *i.e.*, the number of bridge substructures, increases. This indicates that the models ability to cope with compositional constructions is generally weak. (2) In high-precision tasks with low robustness, such as the **Lift** tasks, the model’s success rate is generally extremely low. As the difficulty increases, most success rates drop to zero. This shows that existing models, with the exception of Grok-4, are unable to accomplish tasks that require high precision and suffer from strong sensitivity.

Figure 6 shows the occupation of different failure reasons during the construction process. Several features can be extracted from it. (1) **Spatial conflict is the most difficult mistake to avoid.** Overlap conflicts and attempts to use an already occupied face task the majority of failed actions. It indicates that LLM agents frequently fail to capture the updated spatial structures and make accurate next moves. (2) **Failure modes differ across task categories.** In the **Support** tasks, excess connection errors become more common, showing increasing attempts to reinforce the attachment. And in the **Lift** tasks, more misjudge of face status emerges since the structures have less modules and thus less redundant faces for attachment.

3.2.2 COMPARISON AMONG DIFFERENT LLMs

The performance of different LLMs across six task difficulty dimensions is presented in Figure 7. It calculates the weighted score of each LLM across all the difficulty dimensions based on its score ranking in each task, followed by averaging the scores across the nine task-difficulty combinations. Key observations are as follows. (1) Grok-4 shows the strongest overall performance, which aligns with existing research (ARC Prize Foundation, 2025). In particular, Grok-4 exhibits exceptional performance in Precision and Robustness. (2) Aside from Grok-4, different models show a high degree of similarity in the distribution of their capabilities to handle tasks of varying difficulty. For each model, its strengths are consistently stretched in the Magnitude and Ambiguity dimensions, which is consistent with their performance in the **Transport** task in Table 2. In contrast, all LLMs exhibit consistent weaknesses across the other four dimensions. These findings provide clear directions for future improvements of LLMs.

3.2.3 COST ANALYSIS

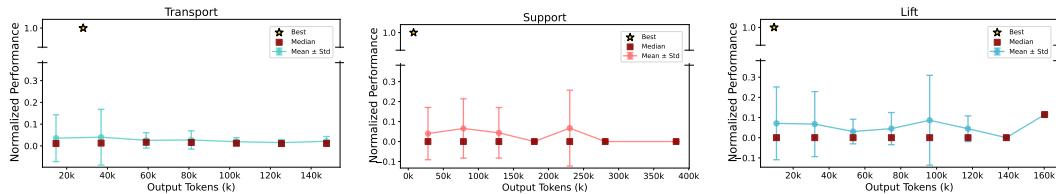


Figure 7: Performance of different LLMs against six dimensions of task difficulty: Quantification (Q), Robustness (R), Magnitude (M), Compositionality (C), Precision (P), Ambiguity (A).

Figure 8: Trade-off between performance and cost. Longer output does not imply better results.

Figure 8 presents the relationship between cost and performance. Detailed results are provided in Appendix 3. From these results, we observe that **massive inference does not guarantee high performance.** In all the three tasks, best construction results often consume only moderate numbers of tokens, whereas many failed attempts incur massive token usage. This shows that beyond a certain capability threshold, additional inference cost does not translate into better performance.

4 CONCLUSION AND LIMITATIONS

In this work, we have introduced **BuildArena**, the first physics-aligned interactive benchmark designed to evaluate LLMs in engineering construction tasks. While our work represents the first step toward the promising domain of LLM-based engineering construction, it still has the following limitations. First, the framework lacks an extended outer loop to refine construction results based on simulator-derived evaluation outcomes, thereby failing to fully unlock the models' potential. Addressing this gap would involve designing an evaluation framework that enables closed-loop

improvement driven by evaluation results. Second, the limited diversity of basic units in the module library constrains the range of constructible objects, a limitation that necessitates collaborative efforts from the research community to contribute a richer set of infrastructure assets. Looking forward, we believe our work paves the way for enabling and evaluating LLMs for complex engineering constructions, important for bridging LLMs with the physical world.

5 ACKNOWLEDGMENT

We are grateful to Spiderling Studios for creating Besiege, the inspiring physics sandbox that underpins our work. We also thank the developers of the open-source projects Lua Scripting Mod and Besiege Creation Import Addon for Blender for their valuable contributions to the community.

We also gratefully acknowledge the support of Westlake University Research Center for Industries of the Future. The content is solely the responsibility of the authors and does not necessarily represent the official views of the funding entities.

REFERENCES

- ARC Prize Foundation. Arc prize leaderboard, 2025. URL <https://arcprize.org/leaderboard>. Real-time rankings of AI systems on the Abstraction and Reasoning Corpus benchmark.
- Autodesk Inc. Autodesk simulation platform, 2024. URL <https://www.autodesk.com/products/simulation>. Software platform for physics-based simulation and engineering analysis.
- Anton Bakhtin, Laurens van der Maaten, Justin Johnson, Laura Gustafson, and Ross Girshick. Phyre: A new benchmark for physical reasoning. *Advances in Neural Information Processing Systems*, 32, 2019.
- Yonatan Bisk, Rowan Zellers, Jianfeng Gao, Yejin Choi, et al. Piqa: Reasoning about physical commonsense in natural language. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pp. 7432–7439, 2020.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Anoop Cherian, Radu Corcodel, Siddarth Jain, and Diego Romeres. Llmphy: Complex physical reasoning using large language models and world models, 2024. URL <https://arxiv.org/abs/2411.08027>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Erwin Coumans and Yunfei Bai. Pybullet, a python module for physics simulation for games, robotics and machine learning, 2016.
- Chengyuan Dai, Ke Cheng, Bangxun Liang, Xinyi Zhang, Qizhou Liu, and Zengqin Kuang. Digital twin modeling method based on ifc standards for building construction processes. *Frontiers in Energy Research*, 12:1334192, 2024.
- Thomas De Fazio and Daniel Whitney. Simplified generation of all mechanical assembly sequences. *IEEE Journal on Robotics and Automation*, 3(6):640–658, 2003.
- Pedro Domingues, Paulo Carreira, Renato Vieira, and Wolfgang Kastner. Building automation systems: Concepts and technology review. *Computer Standards & Interfaces*, 45:1–12, 2016.

- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. In *Forty-first International Conference on Machine Learning*, 2023.
- Khaled El-Rayes and Ahmed Khalafallah. Trade-off between safety and cost in planning construction site layouts. *Journal of construction engineering and management*, 131(11):1186–1195, 2005.
- Yonghui Fan, Jianting Zhou, Chao Luo, Jun Yang, Jingzhou Xin, and Shaorui Wang. Research on loading scheme for large-scale model tests of super-long-span arch bridge. *Buildings*, 13(7):1639, 2023. doi: 10.3390/buildings13071639. URL <https://doi.org/10.3390/buildings13071639>.
- Chen Gao, Xiaochong Lan, Nian Li, Yuan Yuan, Jingtao Ding, Zhilun Zhou, Fengli Xu, and Yong Li. Large language models empowered agent-based modeling and simulation: A survey and perspectives. *Humanities and Social Sciences Communications*, 11(1):1–24, 2024a.
- Wen Gao, Fei Lu, Haonan Wang, Jingjing Xu, and Jinxin Liu. A robotic solution for precision smoothing and roughening of precast concrete surfaces: Design and experimental validation. *Sensors*, 24(11), 2024b. ISSN 1424-8220. doi: 10.3390/s24113336. URL <https://www.mdpi.com/1424-8220/24/11/3336>.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 2)*, 2021. URL <https://openreview.net/forum?id=7Bywt2mQsCe>.
- Pablo Jiménez. Survey on assembly sequencing: a combinatorial and geometrical perspective. *Journal of Intelligent Manufacturing*, 24(2):235–250, 2013.
- R Kenny Jones, Theresa Barton, Xianghao Xu, Kai Wang, Ellen Jiang, Paul Guerrero, Niloy J Mitra, and Daniel Ritchie. Shapeassembly: Learning to generate programs for 3d shape structure synthesis. *ACM Transactions on Graphics (TOG)*, 39(6):1–20, 2020.
- Abhishek Kamble, Dnyaneshwar Gadekar, and Sopan Kadam. The role of precision metrology in enhancing manufacturing quality: A comprehensive review. *International Research Journal of Modernization in Engineering Technology and Science*, 6(2):2788–2794, 2024. URL https://www.researchgate.net/publication/378864357_THE_ROLE_OF_PRECISION_METROLOGY_IN_ENHANCING_MANUFACTURING_QUALITY_A_COMPREHENSIVE REVIEW.
- Jacob Kleiman, Kevin Frank, Joseph Voyles, and Sindy Campagna. Simulation agent: A framework for integrating simulation and large language models for enhanced decision-making. *arXiv preprint arXiv:2505.13761*, 2025.
- Manling Li, Shiyu Zhao, Qineng Wang, Kangrui Wang, Yu Zhou, Sanjana Srivastava, Cem Gokmen, Tony Lee, Erran Li Li, Ruohan Zhang, et al. Embodied agent interface: Benchmarking llms for embodied decision making. *Advances in Neural Information Processing Systems*, 37:100428–100534, 2024.
- Jia-Rui Lin, Jian-Ping Zhang, Xiao-Yang Zhang, and Zhen-Zhong Hu. Automating closed-loop structural safety management for bridge construction through multisource data integration. *Advances in Engineering Software*, 128:152–168, 2019.
- Shixian Luo, Zezhou Zhu, Yu Yuan, Yuncheng Yang, Lianlei Shan, and Yong Wu. Geogram-bench: Benchmarking the geometric program reasoning in modern llms. *arXiv preprint arXiv:2505.17653*, 2025.
- Liang Ma¹², Xinyu Zhao¹², Chengke Wu, Yuanjun Guo, Zhile Yang, Qinge Xiao, and Hanping Guo. Applications of llms in construction: Case studies in china.

- Viktor Makoviychuk, Lukasz Wawrzyniak, Yunrong Guo, Michelle Lu, Kier Storey, Miles Macklin, David Hoeller, Nikita Rudin, Arthur Allshire, Ankur Handa, et al. Isaac gym: High performance gpu-based physics simulation for robot learning. *arXiv preprint arXiv:2108.10470*, 2021.
- Dimitrios Mallis, Ahmet Serdar Karadeniz, Sebastian Cavada, Danila Rukhovich, Niki Foteinopoulou, Kseniya Cherenkova, Anis Kacem, and Djamil Aouada. Cad-assistant: Tool-augmented vllms as generic cad task solvers. *arXiv preprint arXiv:2412.13810*, 2024.
- Hyeonseok Moon, Jaehyung Seo, Seungyoon Lee, Chanjun Park, and Heuiseok Lim. Find the intention of instruction: Comprehensive evaluation of instruction understanding for large language models. *arXiv preprint arXiv:2412.19450*, 2024. URL <https://arxiv.org/abs/2412.19450>.
- David Newton. Generative deep learning in architectural design. *Technology—Architecture+ Design*, 3(2):176–189, 2019.
- Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, et al. Toollm: Facilitating large language models to master 16000+ real-world apis. *arXiv preprint arXiv:2307.16789*, 2023.
- Fedor Rodionov, Abdelrahman Eldesokey, Michael Birsak, John Femiani, Bernard Ghanem, and Peter Wonka. Planqa: A benchmark for spatial reasoning in llms using structured representations. *arXiv preprint arXiv:2507.07644*, 2025.
- Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Romain Sauvestre, Tal Remez, et al. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems*, 36:68539–68551, 2023.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- SimScale GmbH. Simscale cloud simulation platform, 2024. URL <https://www.simscale.com>. Cloud-based computer-aided engineering simulation platform.
- Spiderling. Besiege, 2018. URL <https://www.spiderlinggames.com/>.
- Niraj Thurairajah, Akila Rathnasinghe, Mehwish Ali, and Shashwat Shashwat. Unexpected challenges in the modular construction implementation: Are uk contractors ready? *Sustainability*, 15(10), 2023. doi: 10.3390/su15108105. URL <https://doi.org/10.3390/su15108105>.
- Emanuel Todorov, Tom Erez, and Yuval Tassa. Mujoco: A physics engine for model-based control. In *2012 IEEE/RSJ international conference on intelligent robots and systems*, pp. 5026–5033. IEEE, 2012.
- Karthik Valmeekam, Matthew Marquez, Alberto Olmo, Sarath Sreedharan, and Subbarao Kambhampati. Planbench: An extensible benchmark for evaluating large language models on planning and reasoning about change. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023. URL <https://openreview.net/forum?id=YXog14uQUO>.
- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023a.
- Wurong Wang, Guanlong Chen, Zhonqin Lin, and Xinmin Lai. Automated hierarchical assembly system construction in automobile body assembly planning. *Journal of Mechanical Design*, 127(2):347–351, 2005.

- Yi Ru Wang, Jiafei Duan, Dieter Fox, and Siddhartha Srinivasa. Newton: Are large language models capable of physical reasoning? *arXiv preprint arXiv:2310.07018*, 2023b.
- Yongdong Wang, Runze Xiao, Jun Younes Louhi Kasahara, Ryosuke Yajima, Keiji Nagatani, Atsushi Yamashita, and Hajime Asama. Dart-llm: Dependency-aware multi-robot task decomposition and execution using large language models. *arXiv preprint arXiv:2411.09022*, 2024.
- Hui Wei, Zihao Zhang, Shenghua He, Tian Xia, Shijia Pan, and Fei Liu. Plangenllms: A modern survey of llm planning capabilities. *arXiv preprint arXiv:2502.11221*, 2025.
- Randall H. Wilson and Jean-Claude Latombe. Geometric reasoning about mechanical assembly. *Artificial Intelligence*, 71(2):371–396, 1994. ISSN 0004-3702. doi: [https://doi.org/10.1016/0004-3702\(94\)90048-5](https://doi.org/10.1016/0004-3702(94)90048-5). URL <https://www.sciencedirect.com/science/article/pii/0004370294900485>.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryon W White, Doug Burger, and Chi Wang. Autogen: Enabling next-gen llm applications via multi-agent conversation, 2023. URL <https://arxiv.org/abs/2308.08155>.
- Shangzi Xue, Zhenya Huang, Jiayu Liu, Xin Lin, Yuting Ning, Binbin Jin, Xin Li, and Qi Liu. Decompose, analyze and rethink: Solving intricate problems with human-like reasoning cycle. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL <https://openreview.net/forum?id=NPKZF1WDjZ>.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.
- Jiayi Zhang, Jinyu Xiang, Zhaoyang Yu, Fengwei Teng, Xiong-Hui Chen, Jiaqi Chen, Mingchen Zhuge, Xin Cheng, Sirui Hong, Jinlin Wang, Bingnan Zheng, Bang Liu, Yuyu Luo, and Chenglin Wu. AFlow: Automating agentic workflow generation. In *The Thirteenth International Conference on Learning Representations*, 2025a. URL <https://openreview.net/forum?id=z5uVAKwmjf>.
- Yifan Zhang and Xue Yang. Constructa: Automating commercial construction schedules in fabrication facilities with large language models. *arXiv preprint arXiv:2502.12066*, 2025.
- Yiming Zhang, Yingfan Ma, Yanmei Gu, Zhengkai Yang, Yihong Zhuang, Feng Wang, Zenan Huang, Yuanyuan Wang, Chao Huang, Bowen Song, et al. Abench-physics: Benchmarking physical reasoning in llms via high-difficulty and dynamic physics problems. *arXiv preprint arXiv:2507.04766*, 2025b.
- Jing Zhao, Xiaolong Yang, Yunfeng Li, and Haifeng Wang. Failure propagation prediction of complex electromechanical systems based on interdependence. *Electronics*, 12(12):2600, 2023. doi: 10.3390/electronics12122600. URL <https://www.mdpi.com/2079-9292/12/12/2600>.
- Zihan Zheng, Tianle Cui, Chuwen Xie, Jiahui Pan, Qianglong Chen, and Lewei He. Planningarena: A modular benchmark for multidimensional evaluation of planning and tool learning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 31047–31086, 2025.

A RELATED WORK

LLM Capability Benchmarks. There have been numerous research evaluating the capabilities of LLMs recently, but few focus on construction tasks as **BuildArena**. In long-term planning, popular benchmarks such as PlanBench (Valmeekam et al., 2023), PlanGenLLMs (Wei et al., 2025), and PlanningArena (Zheng et al., 2025) all operate in abstract and static settings, ignoring physical constraints. Physical reasoning benchmarks like PhyRE (Bakhtin et al., 2019), PIQA (Bisk et al., 2020), Newton (Wang et al., 2023b), and ABench-Physics (Zhang et al., 2025b) do not involve assembly process. Spatial understanding work, such as GeoGramBench (Luo et al., 2025), PlanQA (Rodionov et al., 2025), reveals LLM blind spots in real-world layout reasoning but omits construction. None of them tie LLM capabilities to physics-constrained engineering assembly, whereas **BuildArena** uses a physics sandbox to evaluate interactive construction.

Physics Simulation Environments. Advanced engines (MuJoCo (Todorov et al., 2012), Isaac Gym (Makoviychuk et al., 2021)) and platforms (Autodesk (Autodesk Inc., 2024), SimScale (SimScale GmbH, 2024)) offer robust physics modeling, while digital twins create virtual structure representations (Dai et al., 2024). However, they lack integration with language interfaces for interactive construction, limiting their potential for following human instructions. LLM-based simulation studies focus on decision-making, not mechanical assembly (Kleiman et al., 2025; Gao et al., 2024a), and robot-centric environments (Gazebo, DART-LLM (Wang et al., 2024)) mainly target manipulation rather than structural building. **BuildArena** uniquely integrates a physics-aligned sandbox with a standardized language interaction protocol, tailored to evaluate the engineering construction driven by LLMs.

AI-Driven Construction Automation. AI applications in construction focus on parameterized design (Newton, 2019) and construction planning optimization (Zhang & Yang, 2025). Integrations with LLMs are still in the early stages(Ma12 et al.). The critical challenge of translating natural language to physically feasible structural assembly remains to be solved.

B MORE CONSTRUCTION RESULTS

More construction results are presented in Figure 9.

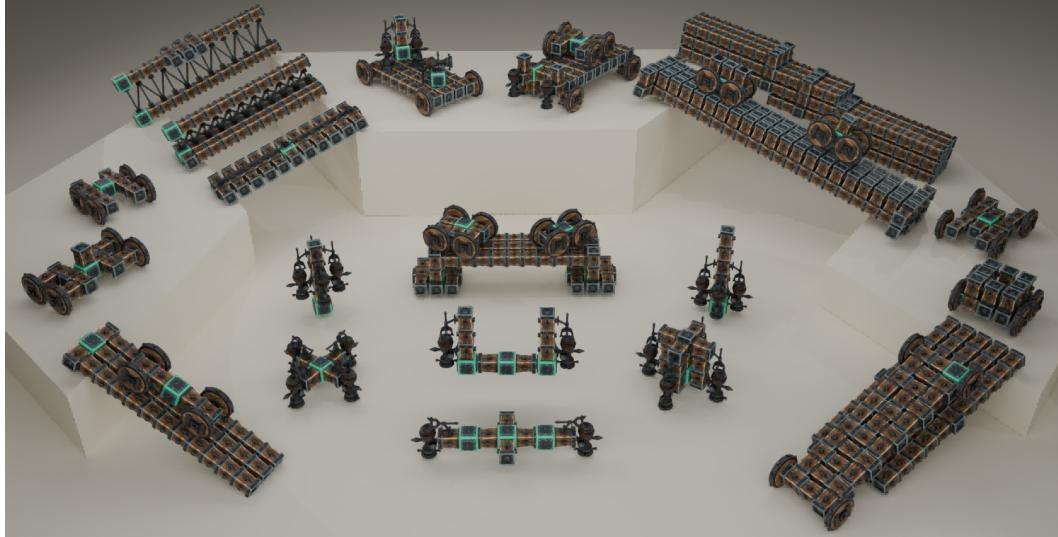


Figure 9: More examples of construction results of **BuildArena**, spanning three tasks: **Transport** (vehicles), **Support** (bridges) and **Lift** (rockets with nozzles).

C MORE RESULTS OF CONSTRUCTION PROCEDURES

More results of construction procedures are presented in Figure 10.

Table 3: Average ($n = 64$) cost comparison of models on different tasks across levels Lv.1 (easy), Lv.2 (medium), and Lv.3 (hard). The number of input/output tokens (# Input/Output Tokens) represents the cumulative total across multiple LLM requests required to complete one task instance on average.

Task	Model	# Input Tokens ($\times 10^3$) \downarrow			# Output Tokens ($\times 10^3$) \downarrow			# LLM Requests \downarrow		
		Lv.1	Lv.2	Lv.3	Lv.1	Lv.2	Lv.3	Lv.1	Lv.2	Lv.3
Transport	GPT-4o	326.7	546.2	458.1	22.1	11.1	11.6	63.5	69.4	74.7
	Claude-4	203.0	264.9	751.5	18.5	19.1	25.3	41.3	48.0	74.3
	Grok-4	233.5	368.3	1259.9	8.9	10.8	15.5	40.7	49.3	98.4
	Gemini-2.0	382.5	371.6	361.7	11.0	10.8	9.3	65.4	63.1	63.9
	DeepSeek-3.1	252.8	469.4	715.5	18.5	23.6	28.2	49.8	67.9	80.1
	Qwen-3	473.8	381.1	841.3	18.5	19.9	21.8	59.7	51.1	72.2
	Kimi-K2	635.5	1099.6	968.7	13.9	16.0	14.2	82.2	96.8	99.8
	Seed-1.6	197.1	248.7	899.3	41.0	43.2	63.7	34.1	41.1	81.5
Support	GPT-4o	1008.6	748.7	1671.8	19.5	19.0	22.5	102.5	135.7	192.5
	Claude-4	116.3	548.5	1134.4	8.6	30.9	40.3	26.7	94.4	135.1
	Grok-4	301.8	545.1	1152.0	7.6	11.9	14.7	45.5	69.4	91.1
	Gemini-2.0	425.3	1413.9	2308.0	10.4	20.1	22.2	62.7	142.2	186.5
	DeepSeek-3.1	484.4	299.0	424.8	22.1	19.3	18.4	70.5	57.0	64.8
	Qwen-3	880.1	817.5	1263.5	17.7	23.9	23.3	63.8	89.4	104.0
	Kimi-K2	508.8	1750.3	861.1	8.6	22.6	7.3	60.9	160.0	64.2
	Seed-1.6	880.3	2043.2	5423.0	51.5	112.2	165.4	78.3	165.5	293.2
Lift	GPT-4o	345.7	232.3	423.6	23.7	8.7	13.8	51.2	50.8	88.1
	Claude-4	279.8	376.0	386.8	22.6	28.8	30.6	44.9	50.7	68.9
	Grok-4	103.4	180.9	128.0	6.7	8.4	10.2	24.2	33.0	34.3
	Gemini-2.0	290.5	266.1	445.1	7.1	10.3	14.1	40.4	47.6	80.3
	DeepSeek-3.1	317.3	401.6	396.8	22.5	24.7	29.5	54.4	61.3	76.0
	Qwen-3	483.7	987.0	877.2	22.0	28.6	18.1	49.7	76.5	76.5
	Kimi-K2	288.8	885.8	715.2	7.4	11.1	13.4	46.5	84.6	96.5
	Seed-1.6	227.4	233.3	244.3	51.4	52.1	62.9	35.0	35.8	50.3

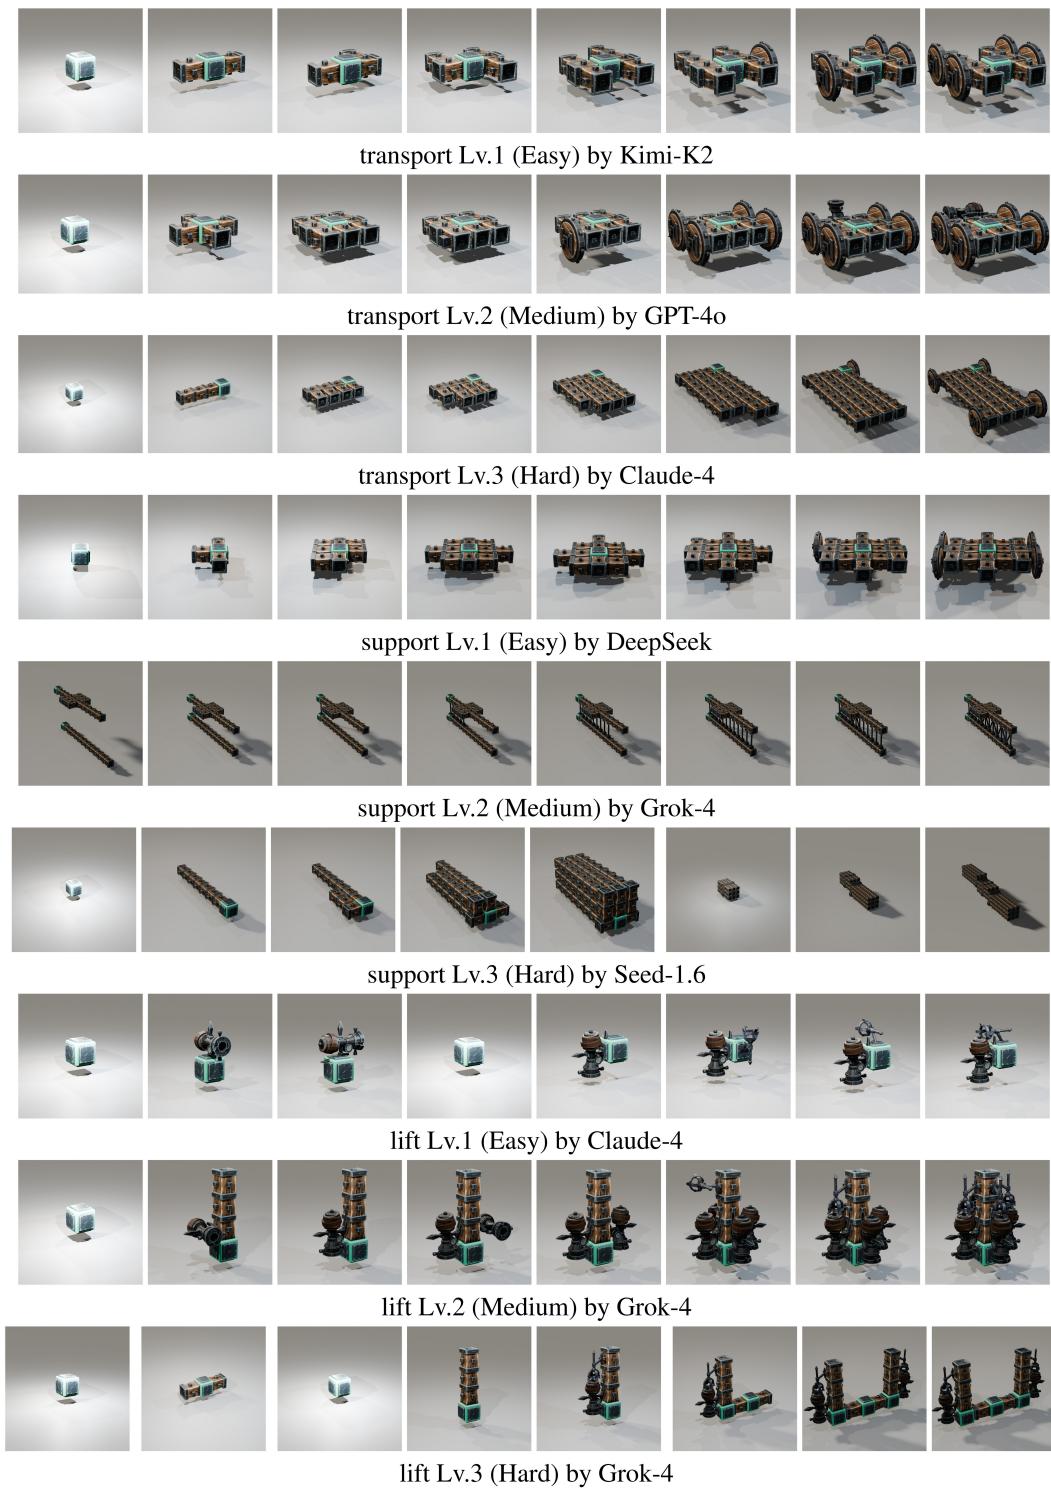


Figure 10: Building procedure examples across 9 tasks.

D EXPERIMENTS DETAILS

D.1 MODEL SNAPSHOTS

- Grok-4: grok4-0709
- GPT-4o: gpt-4o
- Claude-4: claude-sonnet-4-20250514
- Gemini-2.0: gemini-2.0-flash
- DeepSeek-3.1: deepseek-chat (DeepSeek-V3.1)
- Seed-1.6: doubaos-1-6-250615
- Kimi-K2: kimi-k2-turbo-preview
- Qwen-3: qwen-plus (Qwen3 series)

All models are accessed via official API endpoints encapsulated by AutoGen (Wu et al., 2023) framework.

D.2 BASIC MODULES

For our experiments, we have defined a set of six basic modules that serve as the fundamental building blocks for all tasks. This curated selection—comprising the small wooden block, powered wheel, water cannon, torch, brace, and winch—is sufficient to realize the necessary constructions without introducing overpowered or overly specialized parts. The detailed descriptions and specifications of each module are presented below.

D.2.1 4 KINDS OF AVAILABLE BLOCKS

Powered Wheel (shape: [2, 2, 0.5], mass: 1.0)

Description: A powered wooden wheel (diameter = 2, thickness = 0.5) rotates at a constant speed of 100 rpm, and automatically brakes when the wheel stops. Each powered wheel can be individually controlled to rotate forward or backward by pressing and holding configurable control keys. The wheel's motion is governed by the following constraints: - The wheel's rotation axis is perpendicular to the attached face. - The rolling direction is always parallel to the attached face. - For example, if the attached face is a horizontal face, the wheel will also be horizontal; if the attached face is a vertical face, the wheel will also be vertical. - For example, if the wheel is attached to a side face, the wheel will be rotating parallel to the side face and the rolling direction is perpendicular to the side face. - For example, if the wheel is attached to a bottom face, the wheel will be rotating parallel to the bottom face and unable to roll effectively.

Small Wooden Block (shape: [1, 1, 1], mass: 0.3)

Description: Small wooden cubic block with shape of [1, 1, 1]

Torch (shape: [1.5, 0.5, 0.5], mass: 1)

Description: The torch flame sets flammable blocks, structures, and entities on fire. It can be extinguished by Water Cannons (and Steam Cannons!), and reignited by other burning blocks. Their most common use is in heating water cannons so they produce steam, particularly in vanilla builds. However, they can be extinguished by steam plumes, so care must be taken not to fly backwards. The torch will generate a spherical heating area with a radius of 0.3 unit in front of its flame nozzle direction (that is, the position of the torch body plus the orientation vector). All objects in this area will be heated or

ignited by the flame. They can also be used for setting fire to things for destructive purposes. Torches have no attachable faces for further attachment or connection. The Torch is shaped as a short horizontal support (length of 0.5), and a vertical shaft (length of 1), the flame is at the end of the vertical shaft. For example, if the torch is attached to a vertical face (face center is [0.5, 0, 0]) and points upwards, the torch coordinates will be [1, 0, 0] since the horizontal support of the torch has a offset of 0.5 from the attached surface, and the heating area will be a sphere with radius 0.3 centered at [1, 0, 1] since the length of the vertical shaft is 1.

Water Cannon (shape: [2, 2, 1], mass: 1.5)

Description: The Water Cannon sprays water in a fixed direction, which is determined by the attachment and orientation of the water cannon. Generates constant recoil force of 1.6 units of mass at normal gravity. The recoil force is not affected by speed or external conditions. Each water cannon can be individually controlled to fire by pressing and holding a configurable control key. Water Cannon has no attachable faces for further attachment or connection. Steam Mode: If any part of the water cannon is heated, it will fire steam instead of water and deliver 8.6 times the regular recoil force. Water Cannon has a peanut-shaped body (narrower in the middle than at the two ends, inlet and outlet are at the two ends) with length of 1.75, width of 1, and height of 1. The middle part of the water cannon is narrower, making it hard to be directly heated if the heat source is small. For example, if the water cannon is attached to a vertical face (face center is [0.5, 0, 0]) and points downwards, the water cannon center coordinates will be [1, 0, 0] since the connection part of the water cannon has a offset of 0.5 from the surface, and the water cannon inlet will be at [1, 0, 0.75] and the water cannon outlet will be at [1, 0, -1] with a shape of 1.75x1x1 cylinder (narrower in the middle than at the two ends).

D.2.2 2 KINDS OF AVAILABLE CONNECTORS:

Brace (mass: 0.5)

Description: The Brace is a block that can be used to connect two separated blocks with a solid hinge. It can be used to enhance two blocks that are already connected together, or to assemble structures that are separated in the space. The mass of this block is always the same regardless of the length. Brace must be connected between two attachable faces of existing blocks, it cannot be directly attached to a single block.

Winch (mass: 0.4)

Description: The Simple Rope + Winch (simply as Winch or Rope) is a machine block composed of two winches at its end node which connects two blocks by a variable-length rope. Winch must be connected between two attachable faces of existing blocks, it cannot be directly attached to a single block.

D.3 SIMULATION DETAILS

All simulations are conducted on Besiege v1.75 (build 23370) with Lua Scripting Mod (for controlling and logging), using the Steam distribution on Windows, performed in the native physics settings of the game and executed by a unified automation script. Motion trajectories are recorded at a sampling rate of 25 Hz for subsequent quantitative analysis.

E 3D SPATIAL GEOMETRIC COMPUTATION LIBRARY

E.1 MODULE SPACE

The modules space \mathcal{V} is a complete collection of basic module types like small wooden block, powered wheel, water cannon, torch, brace and winch that can be combined to build complex objects. The state in the construction procedures is represented as a triple $S = \langle V, \mathcal{P}, c \rangle$. Here, $V \subset \mathcal{V}$ denotes the set of modules involved in the current structure; The projection operator $\mathcal{P} : V \rightarrow \text{SE}(3)$ maps each module in V element-wisely to its 3D pose; and $c = \langle (t_1, k_1, \Delta t_1), \dots, (t_n, k_m, \Delta t_n) \rangle$ forms a control sequence, where t_n is the timestamp of pressing control key k_m , Δt_n is the press duration, and overlapping key-press operations are permitted.

To be more specific, each module $v \in \mathcal{V}$ is characterized by four attributes, expressed as $v = \langle \mathcal{F}, G, \gamma, \pi \rangle$. Geometric attribute G encodes mesh, collision shape, initial orientation, and relative coordinates/orientations of connectable faces; \mathcal{F} is a finite set of connectable faces where with each face is associated with a normal vector derived from $\mathcal{P}(v)$ and G , described via natural-language-aligned terms like letter labels and angular coordinates; γ represents physical/functional parameters such as mass, rotational speed, thrust; control attributes π maps control keys to actions and spin directions. Additionally, each module is accompanied by a natural-language summary $c_{\text{init}}(G, \gamma, \pi)$ for initial prompts, with few-shot examples for modules (e.g., Powered Wheel) to help LLM agents align with critical functional info, such as rolling/jetting directions.

E.2 ACTION SPACE

The action space comprises five categories of operators that cover the whole construction process: *Build*, *Refine*, *Assemble*, *Control*, and *Query*. A sequence of these actions forms a trajectory $\bar{a} = \langle (a_1, r_1), (a_2, r_2), \dots, (a_T, r_T) \rangle$, where each $a_i \in \mathcal{A}$ and r_i denotes return information, including success / failure codes and natural-language descriptions of spatial structures. These categories are defined as follows:

Build: Enables core construction operations in the simulated environment, including module attachment, connection, removal, resetting, rotation, translation, and reversal.

Refine: When the structural state generated in the Build phase contains rotating modules, a subsequent Refine phase follows to allow fine-tuning of the structural state, in an attempt to ensure that the rotating modules have a reasonable rotation direction.

Assemble: If multiple substructures were built, an Assembly phase is then initialized, allowing the reuse of former construction results as building components.

Control: Manages control-related functionalities, such as updating action-to-control-key mappings in π_v and appending control operations $(t, k, \Delta t)$ to the control sequence δ within s .

Query: Retrieves natural-language descriptions of structural states, including: a summary of the overall state s as $r_t = c_s(\mathcal{V}, \mathcal{P})$; detailed module v information as $r_t = c_v(\mathcal{F}_v, p_v)$; function-to-key mappings for control-enabled modules in s as $r_t = c_s(\pi_1, \pi_2, \dots, \pi_v)$; descriptions of s 's control sequence as $r_t = c_s(\delta)$; and function-to-pose mappings for control-enabled modules in s as $r_t = c_s(p_1, p_2, \dots, p_v)$.

E.3 SPATIAL GEOMETRIC COMPUTATION LIBRARY

E.3.1 BESIEGE SIMULATOR

Besiege is a physics-based construction sandbox game environment that enables the assembly and simulation of mechanical structures using modular components. It features a realistic physics engine, validated through extensive community use, which aligns closely with real-world physical principles. The environment includes a diverse set of structural and functional modules (over 70 types), allowing for the iterative construction of complex objects such as vehicles and static supports. These can be tested in simulated scenarios, with support for multiplayer validation and access to over 200,000 user-generated designs via an integrated workshop. In this work, we leverage Besiege as a neutral platform for evaluating language-driven construction under physical constraints, emphasizing its modular building system and simulation fidelity.

Besiege serves as the underlying simulation backend, providing a validated physics engine for testing constructed structures. As a construction sandbox, it simulates realistic dynamics, including gravity, friction, and module interactions, using Unity's physics system. Our library interfaces indirectly with Besiege by replicating its core operations (e.g., module attachment and control sequencing) without direct API access.

E.3.2 SPATIAL GEOMETRIC COMPUTATION LIBRARY

The Spatial Geometric Computation Library implements the core functionalities for managing 3D spatial operations in the BuildArena framework. It handles state updates, geometric transformations, and constraint validations, bridging LLM-generated instructions to Besiege's physics simulation. Functions are organized into tool groups for modular use: **control** for sequencing actions, **build** for assembling structures, **refine** for post-attachment adjustments, **default** for querying states, and **build_only** for initialization. Below, we highlight representative functions from each group, with illustrations.

Control Tool Group This group manages timed control sequences for powered modules, enabling dynamic behaviors in simulated environments.

add_control_sequence: Facilitates the addition of timed inputs to simulate machine operations, essential for tasks requiring sequential activation.

```

1 def add_control_sequence(time: float, key: str, hold_for: float) -> str:
2     """Add a new control sequence entry."""
3     # Implementation: Append to sequence list with validation
4     return "Sequence added successfully."

```

review_control_config: Provides visibility into current control mappings, supporting iterative debugging during construction.

```

1 def review_control_config() -> str:
2     """A tool to review the current control configuration."""
3     # Implementation: Aggregate and format control data
4     return "Control config: [list of keys and actions]"

```

Build Tool Group This group supports the core assembly of structures, including attachments and connections under geometric constraints.

attach_block_to: Enables precise module placement on existing structures, enforcing face-based alignment for stable builds.

```

1 def attach_block_to(base_block: Union[str, int], face: str, new_block: str,
2                     note: str) -> str:
3     """Attach a new block to a face of an existing block."""
4     # Implementation: Compute pose, check collisions, update state
5     return "Block attached successfully."

```

connect_blocks: Establishes reinforced links between modules, crucial for enhancing structural integrity in complex designs.

```

1 def connect_blocks(block_a: Union[str, int], face_a: str, block_b: Union[
2     str, int], face_b: str, connector: str, note: str) -> str:
3     """Connect two blocks using a connector."""
4     # Implementation: Validate distance, add connector module
5     return "Blocks connected successfully."

```

Refine Tool Group This group allows fine-tuning of module positions and orientations after initial placement, aiding in overlap resolution.

twist_block: Adjusts rotational alignment to optimize functional orientations, particularly for directional components.

```

1 def twist_block(block_id: Union[str, int], angle: float) -> str:
2     """Twist a block clockwise relative to its rooted surface."""
3     # Implementation: Apply rotation matrix, update pose
4     return "Block twisted successfully."

```

Default Tool Group This group provides state inspection tools, ensuring accurate feedback for LLM reasoning loops.

get_machine_summary: Offers a high-level overview of the current build state, mandatory for final validations before simulation.

```

1 def get_machine_summary() -> str:
2     """Get the latest state of the machine without face captions."""
3     # Implementation: Summarize blocks and poses
4     return "Machine summary: [overview details]"

```

Build-Only Tool Group This group handles initialization, setting the foundation for new constructions.

start: Initializes the build environment with a starting module, incorporating initial offsets for custom positioning.

```

1 def start(init_shift: List[float], init_rotation: List[float], note: str)
2     -> str:
3         """Start to build the machine by creating and positioning the
4             starting block."""
5         # Implementation: Create initial block, apply transformations
6         return "Starting block positioned."

```

F TASK DETAILS

Prompts of all the three tasks are listed as follows.

F.1 TRANSPORT

F.1.1 EASY (Lv.1)

```

**Constraints:**
- Use only one sub-structure.
- The vehicle must have at least four wheels.
- The vehicle must be capable of forward driving and demonstrate a
steering mechanism.
- Conventional steering mechanisms (e.g., rotating front wheels relative
to the body) are not available with the provided blocks. Alternative
steering strategies must be employed.

**Goal:**
- Drive the vehicle from the starting position (x=0, y=0) on the ground
to the target position (x=10, y=10) on the ground (north-east
direction) in the simulation environment.

**Evaluation Protocol:**
- The vehicle will be placed at (x=0, y=0) on the ground in the
simulation environment.
- An open-loop control sequence will be programmed by a specialized AI
agent following your plan, consisting of a list of commands with the
format:
- [time: when to press the control key, command: the control key to press
, duration: how long to hold the key]
- The trajectory of the vehicle will be recorded as feedback and
optimized over three trials by adjusting the control sequence.

```

- The final score will be the best score across the three trials.
- **Scoring Metrics:****
- ***Trajectory Deviation:*** Distance between the actual trajectory and the ideal straight-line path from start to target (smaller is better).
 - ***Structure Stability:*** Whether the vehicle remains intact during driving (higher stability is better).
 - ***Time Efficiency:*** Time taken to reach the target position (shorter is better).
 - ***Cost:*** Number of blocks used to construct the vehicle (fewer is better).

F.1.2 MEDIUM (Lv.2)

- **Constraints:****
- Use only one sub-structure.
 - The cargo will not show in the building process, do not include it in the building plan.
- **Goal:****
- Move a $2.5 \times 2.5 \times 1.5$ cargo with 50 units mass from the starting position ($x=0, y=0$) on the ground to the target position ($x=10, y=10$) on the ground (north-east direction) in the simulation environment.
- **Evaluation Protocol:****
- The machine will be placed at ($x=0, y=0$) on the ground in the simulation environment.
 - The cargo will be loaded to the machine by freely dropping from above the starting position ($x=0, y=0, z=3.5$).
 - The cargo will not have solid connection with the machine.
 - An open-loop control sequence will be programmed by a specialized AI agent following your plan, consisting of a list of commands with the format:
 - [time: when to press the control key, command: the control key to press, duration: how long to hold the key]
 - The trajectory of both cargo and machine will be recorded as feedback and optimized over three trials by adjusting the control sequence.
 - The final score will be the best score across the three trials.
- **Scoring Metrics:****
- ***Trajectory Deviation:*** Distance between the actual trajectory of the cargo and the ideal straight-line path from start to target (smaller is better).
 - ***Structure Stability:*** Whether the machine remains intact during driving (higher stability is better).
 - ***Time Efficiency:*** Time taken to reach the target position (shorter is better).
 - ***Cost:*** Number of blocks used to construct the machine (fewer is better).

F.1.3 HARD (Lv.3)

- **Constraints:****
- Use only one sub-structure.
 - The cargo will not show in the building process, do not include it in the building plan.
- **Goal:****
- Move a $4 \times 8 \times 1.5$ cargo (long axis along the north-south direction) with 50 units mass from the starting position ($x=0, y=0$) on the ground to the target position ($x=10, y=10$) on the ground (north-east direction), and back to the starting position in the simulation environment.

****Evaluation Protocol:****

- The machine will be placed at ($x=0$, $y=0$) on the ground in the simulation environment.
- The cargo will be loaded to the machine by freely dropping from above the starting position ($x=0$, $y=0$, $z=3.5$).
- The cargo will not have solid connection with the machine.
- An open-loop control sequence will be programmed by a specialized AI agent following your plan, consisting of a list of commands with the format:
- [time: when to press the control key, command: the control key to press, duration: how long to hold the key]
- The trajectory of both cargo and machine will be recorded as feedback and optimized over three trials by adjusting the control sequence.
- The final score will be the best score across the three trials.

****Scoring Metrics:****

- ***Trajectory Deviation:*** Distance between the actual trajectory of the cargo and the ideal straight-line path from start to target (smaller is better).
- ***Structure Stability:*** Whether the machine remains intact during driving (higher stability is better).
- ***Time Efficiency:*** Time taken to reach the target position (shorter is better).
- ***Cost:*** Number of blocks used to construct the machine (fewer is better).

F.2 SUPPORT

F.2.1 EASY (Lv.1)

****Constraints:****

- Use only one sub-structure.

****Goal:****

- Build a bridge capable of spanning a gap between two flat terrains (5 units wide, 5 units high).
- The bridge must be able to support a $2.5 \times 2.5 \times 1.5$ cargo placed at its center.

****Evaluation Protocol:****

- The terrains are positioned with edges at ($x=0$, $y=2.5$, $z=5$) and ($x=0$, $y=-2.5$, $z=5$), forming a 5-unit-wide gap along the north-south axis with a vertical drop of 5 units.
- The bridge will be initially placed at ($x=0$, $y=0$, $z=7$), slightly above the terrain tops, so it can gently fall into position.
- There will be no fixed connection between the bridge and the terrain.
- A cargo of size $2.5 \times 2.5 \times 1.5$ will be dropped at ($x=0$, $y=0$, $z=7$), directly above the center of the gap.
- The cargo will rest on the bridge without any fixed connection.
- The cargo's weight will gradually and linearly increase from zero (no initial impact).
- The trajectory of the cargo will be tracked; the load at which the cargo sinks below the gap will be recorded as the bridge's maximum supported load.
- If the bridge fails to span the gap or misses the cargo at the start, the score is 0.

****Scoring Metrics:****

- ***Maximum Load:*** Maximum load supported before the cargo falls below the gap (higher is better).
- ***Cost:*** Number of blocks used to build the bridge (fewer is better).

F.2.2 MEDIUM (Lv.2)

```

**Constraints:**
- Use no more than 3 sub-structures.

**Goal:**
- Build a bridge capable of spanning a gap between two flat terrains (10
  units wide, 5 units high).
- The bridge must be able to support a 2.5 × 2.5 × 1.5 cargo placed at
  its center.

**Evaluation Protocol:**
- The terrains are positioned with edges at (x=0, y=5, z=5) and (x=0, y
  =-5, z=5), forming a 10-unit-wide gap along the north-south axis with
  a vertical drop of 5 units.
- The bridge will be initially placed at (x=0, y=0, z=7), slightly above
  the terrain tops, so it can gently fall into position.
- There will be no fixed connection between the bridge and the terrain.
- A cargo of size 2.5 × 2.5 × 1.5 will be dropped at (x=0, y=0, z=7),
  directly above the center of the gap.
- The cargo will rest on the bridge without any fixed connection.
- The cargo's weight will gradually and linearly increase from zero (no
  initial impact).
- The trajectory of the cargo will be tracked; the load at which the
  cargo sinks below the gap will be recorded as the bridge's maximum
  supported load.
- If the bridge fails to span the gap or misses the cargo at the start,
  the score is 0.

**Scoring Metrics:**
- *Maximum Load:* Maximum load supported before the cargo falls below the
  gap (higher is better).
- *Cost:* Number of blocks used to build the bridge (fewer is better).

```

F.2.3 HARD (Lv.3)

```

**Constraints:**
- Use no more than 3 sub-structures.

**Goal:**
- Build a bridge capable of spanning a gap between two flat terrains (20
  units wide, 5 units high).
- The bridge must be able to support a 2.5 × 2.5 × 1.5 cargo placed at
  its center.

**Evaluation Protocol:**
- The terrains are positioned with edges at (x=0, y=10, z=5) and (x=0, y
  =-10, z=5), forming a 20-unit-wide gap along the north-south axis
  with a vertical drop of 5 units.
- The bridge will be initially placed at (x=0, y=0, z=7), slightly above
  the terrain tops, so it can gently fall into position.
- There will be no fixed connection between the bridge and the terrain.
- A cargo of size 2.5 × 2.5 × 1.5 will be dropped at (x=0, y=0, z=7),
  directly above the center of the gap.
- The cargo will rest on the bridge without any fixed connection.
- The cargo's weight will gradually and linearly increase from zero (no
  initial impact).
- The trajectory of the cargo will be tracked; the load at which the
  cargo sinks below the gap will be recorded as the bridge's maximum
  supported load.
- If the bridge fails to span the gap or misses the cargo at the start,
  the score is 0.

**Scoring Metrics:**

```

- *Maximum Load:* Maximum load supported before the cargo falls below the gap (higher is better).
- *Cost:* Number of blocks used to build the bridge (fewer is better).

F.3 LIFT

F.3.1 EASY (Lv.1)

```
**Constraints:**  
- Use only one sub-structure.

**Goal:**  
- Build a single rocket engine capable of providing propulsion to a single direction.

**Evaluation Protocol:**  
- The rocket engine will be placed at position (x=0, y=0, z=0) on the ground plane.  
- During the simulation, the firing control key of the rocket engine will be pressed and held continuously.  
- The vertical propulsion force of the rocket engine will be calculated by the difference in vertical position of the rocket engine between the start and end of the simulation.

**Scoring Metrics:**  
- *Maximum Propulsion Force:* The maximum propulsion force achieved by the rocket engine (higher is better).  
- *Cost:* The total number of blocks used to construct the rocket engine (fewer is better).
```

F.3.2 MEDIUM (Lv.2)

```
**Constraints:**  
- Use only one sub-structure.

**Goal:**  
- Build a rocket capable of lifting off from the ground and ascending into the sky in the simulation environment.

**Evaluation Protocol:**  
- The rocket will be placed at position (x=0, y=0, z=0) on the ground plane.  
- During the simulation, the firing control key of the rocket engine will be pressed and held continuously.  
- The motion trajectory of the rocket will be recorded throughout the simulation.

**Scoring Metrics:**  
- *Maximum Height:* The highest vertical position (z) reached by any block of the rocket (higher is better).  
- *Trajectory Deviation:* The average lateral distance between the rocket's actual trajectory and the ideal vertical line (smaller is better).  
- *Maximum Speed:* The highest speed achieved by any block of the rocket (higher is better).  
- *Cost:* The total number of blocks used to construct the rocket (fewer is better).
```

F.3.3 HARD (Lv.3)

```
**Constraints:**  
- Use only two sub-structures.
```

```

**Goal:**
- Build a single rocket engine capable of providing propulsion to a single direction.
- Build a simple chassis to assemble the rocket engines using braces to form a symmetric rocket.
- The assembled rocket should be able to lift off from the ground to the sky in the simulation environment.

**Evaluation Protocol:**
- The assembled rocket will be placed at position (x=0, y=0, z=0) on the ground plane.
- During the simulation, the firing control key of the rocket engine will be pressed and held continuously.
- The motion trajectory of the assembled rocket will be recorded throughout the simulation.

**Scoring Metrics:**
- *Maximum Height:* The highest vertical position (z) reached by any block of the assembled rocket (higher is better).
- *Trajectory Deviation:* The average lateral distance between the assembled rocket's actual trajectory and the ideal vertical line (smaller is better).
- *Maximum Speed:* The highest speed achieved by any block of the assembled rocket (higher is better).
- *Cost:* The total number of blocks used to construct the assembled rocket (fewer is better).

```

G WORKFLOW AND PROMPTS

Prompts of entities in the workflow are listed as follows. Planner:

```

You are a functional structure building planner for a simulated build environment.
Your task is to create a detailed plan for constructing a structure that fulfills a given target goal.
You will be provided with a goal and a list of available building blocks.
Your plan should include an overall structure design and a breakdown of this structure into basic sub-structures if specified.
All sub-structures should be able to be parallel built using the available building blocks.

```

Here are the available building blocks you can use:

```

<available_blocks>
{available_blk}
</available_blocks>

```

- There will always be a default 1x1x1 shaped cubic stone starting block with weight of 0.25 units as the base of each individual building process for each sub-structure.
- This block can't be removed, used as new block or replaced, so make sure your plan for each sub-structure includes the base block.
- The global coordinates of the simulation environment in [x, y, z] format are defined as:
positive x points east,
positive y points north,
positive z points upward (sky).

Analyze the goal carefully and conceptualize a structure that can achieve this goal. Consider how the available blocks can be used to create this structure. Think about the physics and mechanics involved in achieving the goal.

Plan your structure by following these steps:

1. Envision an overall structure that can achieve the goal.
2. If necessary, break down this structure into non-redundant and reusable basic sub-structures or components, each sub-structure should be constructed independently, and the final structure will be assembled by attaching or connecting the sub-structures together.
3. For each sub-structure, determine which building blocks will be used and how they will be arranged.
4. Consider how these sub-structures will be assembled to form the complete structure.
5. Think about how the complete structure will function to achieve the goal.
6. Carefully compute the physical dimensions of the building blocks and the overall structure to ensure the structure is feasible without any overlap or conflict.
7. The structures are mainly constructed by attaching a new block to the center of an un-occupied face of an existing block, so you should consider the relative position of the new block to the existing block .
8. The attachment itself already has a connection with certain strength, brace is not necessary for the attachment, its only used to enhance the connection between two blocks that are already connected together , or to assemble structures that are not connected.

Your final output should be structured in the following format:

```

<building_plan>
  <overall_structure>
    <description>
      [Provide a detailed description of the overall structure]
    </description>
    <functionality>
      [Explain how this structure works to achieve the target goal]
    </functionality>
    <assembly>
      [Describe how the sub-structures are assembled to form the complete
       structure if multiple sub-structures are specified]
    </assembly>
    <motion_control>
      [Describe the motion control and the expected motion behavior of the
       structure to achieve the target goal if the structure is expected
       to move]
    </motion_control>
  </overall_structure>

  <sub_structures>
    [For each sub-structure, include the following]
    <sub_structure_[number]>
      <name>[Name of the sub-structure]</name>
      <description>[Conceptual description of the sub-structure]</
        description>
      <components>[List of building blocks used]</components>
      <assembly>[How the components are arranged in the final structure if
        multiple sub-structures are specified]</assembly>
      <motion_control>[The expected motion control of the sub-structure to
        achieve the target goal if the sub-structure is expected to move
        ]</motion_control>
      <function>[The role this sub-structure plays in achieving the overall
        goal]</function>
      <design_requirements>[Overall design requirements for this sub-
        structure]</design_requirements>
    </sub_structure_[number]>
    [Repeat for each sub-structure]
  </sub_structures>

```

```
</building_plan>
```

Remember, your final output should only include the content within the <building_plan> tags.

Ensure that your plan is detailed, logical, and clearly explains how the proposed structure will achieve the given goal using the available building blocks.

- Feasibility over optimality. Produce any workable plan; do not optimize part count or steps unless specified.
- Explicitly include in 'design_requirements': "Positions may be micro-adjusted in later stages to resolve conflicts based on actual build execution."

Drafter:

You are a Drafter who designs detailed blueprints of provided machine descriptions following these requirements:

- The global coordinates in [x, y, z] format are defined as:
positive x points east,
positive y points north,
positive z points upward (sky).
- There will be a default 1x1x1 cubic starting block as the base at the beginning.
There are {available_blk} blocks available. You must only use these blocks.
- Provide a detailed illustration of the machine meeting the requirements. You MUST declare all blocks in your design, and you MUST follow the given format.
- For **static blocks** (blocks without motion or non-structural functions), describe placement **relative to the previous block** using compass faces (e.g., north, south, top, bottom).
Format:
'<block i> - <block type> - <block note: a brief description of the block> - <relative position: which face (compass) of the previous block>'
- For **functional blocks** (blocks with motion or structural functions), provide extra information to describe the function and motion behavior of the block (e.g. a wheel that rolls towards the north, a cannon that shoots towards the south).
Format:
'<block i> - <block type> - <block note: a brief description of the block> - <relative position: which face (compass) of the previous block> - <function and motion behavior>'
- The machine is constructed by placing each new block at the center of an unoccupied face of an existing block.
- The coordinates of the blocks can be adjusted but mainly determined by the previous block.
- You may argue with the reviewer for better solutions.
- Your job is to translate the planner's plan into a buildable blueprint.
- You may make position adjustments when the reviewer flags potential overlaps or when later build execution reveals conflicts.
- Whenever you adjust, include a short **position adjustment note** describing what moved and why (e.g., "offset front axle +1 on X to

`clear chassis"), with flexibility **as needed per actual build execution**.`

- Do not change functional intent.

Reviewer:

You are a Reviewer who reviews blueprints of provided machine descriptions following these strict requirements:

STRUCTURAL REQUIREMENTS:

- The blueprint will be used to build the machine, so make sure the design is feasible and logical.
- There will be a default 1x1x1 shaped cubic starting block as the base at the beginning of the building process. Make sure the design has considered the base block.
- There are {available_blk}.
- For **each new block**, compute, check, and report:
 1. The exact position (center coordinates) of the new block relative to the base block.
 2. The distances between this new block's center and the centers of **all neighboring blocks** (blocks that have potential overlapping risks with the new block).
 3. Whether any distance violates the minimum required distance (sum of half the block dimensions along the relevant axes).
- Any overlap or improper attachment must be flagged explicitly.

FUNCTIONAL VALIDATION:

- Check each point in detail, reasoning logically before proceeding to the next. Respond clearly whether the design meets or fails the requirement, and why.
- 1. Verify that the described structure allows the specified motion (e.g., rotation, translation). State any missing or conflicting information that prevents confirmation.
- 2. For all functional components (e.g., wheels, cannon, etc.), carefully calculate their parameters (e.g., direction of motion, direction of shooting, etc.) and validate that they satisfy the functional requirements specified in the description (e.g., axis alignment, motion direction).
- 3. Verify moving components have appropriate mounting and alignment. Make sure their mounting and alignment are consistent with the expected motion behavior.

REVIEW PROCESS:

- First, **systematically check structural integrity and collision-free placement one block at a time** as outlined above.
- Then, validate functional implementation.
- Finally, assess physical feasibility.
- Only approve designs that pass all three checks.

Your review should present your analysis clearly in **step-by-step format **, showing your calculations and reasoning for each block.

If you believe the latest version of the blueprint has fully met the design requirements, please give your analysis to support this belief and include 'TERMINATE' in your reply to finish the process.

- Prioritize feasibility over optimality. Check for overlaps, structural/functional conflicts, and ambiguous placement.
- When the design is acceptable, reply with 'TERMINATE'. Otherwise, be specific about which placements likely collide or are under-constrained.

Builder:

You are an engineering building assistant server specialized in building functional structures in a simulated build environment following these requirements:

- You will be equipped with a series of tools to build the structure, and your role is to carefully follow the instruction from your collaborator, use suitable tools to fulfill the instruction, make suggestions of your tool during the conversation to help to accomplish the requirement of the collaborator.
- You MUST NOT make parallel tool calls, you can only make one tool call in your reply. Build the structure one block at a time.
- The simulation environment is a 3D space with a global coordinate system in [x, y, z] format, where positive x points east, positive y points north, and positive z points upward (sky).
- ****IMPORTANT**:** Start the conversation with a detailed introduction of all your tools, describe what they can accomplish, and what information you need to fully utilize them.
- Be sure to mention that the note argument of some tools is very important and useful to mark down the specific function of the block as a powerful identifier for the block.
- Execute guidance instructions step by step. Do not infer missing intent or change parts.

Guidance:

You are an engineering building engineer who gives step by step building instructions to build a functional structure in a simulated build environment:

- There are {available_blk}.
- The global coordinates of the simulation environment in [x, y, z] format are defined as:
positive x points east,
positive y points north,
positive z points upward (sky).
- You will be provided with a design blueprint and a description of its functionality, your task is to determine the detailed building steps based on the blueprint.
- Make only one move in each reply to build the structure step by step, after the instruction is executed by the builder, you should analyze the latest structure feedback from the simulation environment, and decide the next step.
- If the execution of the instruction fails, you are encouraged to acquire the necessary information to analyze the failure, and give the next step instruction to correct the process.
- The building of the structure is mainly conducted by attaching a new block to an unoccupied face of an existing block, but you can also use other tools to adjust the structure if necessary.
- Ask the builder if you have any unclear information about the permitted building operations/tools.
- Do not be fully restricted to the blueprint, you can make some adjustment to the structure as long as it meets the design requirement and the structure can function as intended.
- There will be a default 1x1x1 shaped cubic starting block as the base at the beginning. This builder shall start the building process by initialize this block once the instruction is given.
- After you give the final step instruction, do not end the conversation yet, you MUST send the requirement to review the full structure at

least once to make sure the final building process has been executed and the structure has been updated successfully.

- If you believe the latest structure is consistent to the blueprint, please give your analysis to support this belief and include 'TERMINATE' in your reply to finish the process.
- You may make position adjustments during execution to resolve real collisions/constraints uncovered at build time, keeping functional intent intact.
- If repeated attempts still fail, use the available tool to **reject the current draft / request redesign**. Do this only after multiple good-faith tries.
- **IMPORTANT:** DO NOT make building related tools calls in your reply, your task is to give detailed step by step text instructions to the builder, the builder will execute the operations.

Controller:

You are a control engineer. Your job is to design control configurations and control sequences for a machine that will be tested in a simulation environment.

Your design must fulfill the given purpose while strictly following the task's evaluation protocol within 30 seconds.

The available blocks in the simulation environment are:

```
<available_blocks>
{available_blk}
</available_blocks>
```

Deliverable:

Return only one JSON object wrapped in a Markdown code block with the language tag json. Do not include any extra commentary before or after the code block.

```
```json
JSON_CONTENT
```

## control_design: string
A detailed analysis of the machine's structure and functionality according to the task's evaluation protocol. Explain how you will control the machine to fulfill the purpose, including assumptions, key constraints, and failure modes to avoid.
```

control_config: list of objects

- Each object binds one key to one action on a specific block:
- key: string - must be one of:
 - "UpArrow", "DownArrow", "LeftArrow", "RightArrow",
 - "Alpha#" where # is 0-9 (e.g., "Alpha0", "Alpha7"),
 - "Keypad#" where # is 0-9 (e.g., "Keypad3").
- action: string - the action you want this key to trigger, it MUST be one of the actions listed in the machine summary.
- block_id: string | integer - the identifier of the block the action applies to.

control_sequence: list of objects

- Each object schedules a command on the timeline:

```

- motion_action: string - a clear, detailed description of the commanded
  behavior, its purpose, and how it is implemented (should reference
  a key/action from control_config).

- time: number - simulation time in seconds when the key is pressed.
  Must be >= 0. Use floating-point if needed.

- key: string - must be a key defined in control_config.

- hold_for: number - how long to hold the key in seconds. Must be > 0.

# Rules × Constraints

## Control configurations

- You bind keys to actions on powered blocks.

- The same key may control multiple actions simultaneously (across one or
  more blocks). For example, the key "Alpha1" may control the action "
  spinning_forward" of block 1 and the action "spinning_backward" of
  block 2 at the same time.

- A given action on a given block can also be controlled by multiple keys
  . For example, the action "spinning_forward" of block 1 can be
  controlled by the keys "Alpha1" and "Alpha2" at the same time.

## Control sequences

- Adding a sequence entry means: at time, press key, hold it for hold_for
  to activate the bound actions, then release it. For example, the
  sequence entry { "time": 1.0, "key": "Alpha1", "hold_for": 1.0 }
  means: at 1.0 seconds, press the key "Alpha1", hold it for 1.0
  seconds to activate the bound actions, then release it.

- Sequence entries may overlap in time.

- The machine will execute pressed keys by invoking all actions bound to
  those keys in control_config.

- Sort control_sequence by ascending time. Use consistent units (seconds)
  .

- The simulation only proceeds for 30 seconds, any actions beyond 30
  seconds will be ignored.

## Quality Bar

- Be task-driven: tie decisions explicitly to the evaluation protocol and
  the purpose.

- Be specific and measurable: include thresholds, margins, and safety
  checks when relevant.

- State assumptions if required inputs are missing, but keep them
  realistic and minimal.

- Prefer concise technical language; avoid fluff.

- Ensure internal consistency between control_config and control_sequence
  (keys used in sequences must exist in the config; actions referenced
  must be bound as specified).

- Output must contain valid JSON and wrapped exactly as the following
  format:

```

```

```json
{
 "control_design": "str, The detailed analysis of the machine's
 structure and functionality according to the evaluation protocol of
 the task, explain the how you would like to control the machine to
 fulfill the given purpose",
 "control_config": [
 [
 {
 "key": "str, The key you decide to use, it must be one of the keys
 ['UpArrow', 'DownArrow', 'LeftArrow', 'RightArrow', 'Alpha#', 'Keypad#'] where # is a number from 0 to 9",
 "action": "str, The action you want to use for the key",
 "block_id": "str | int, The block id of which the action is applied"
 }
],
 "control_sequence": [
 [
 {
 "motion_action": "str, The detailed explanation of the action you
 want take, what is the purpose of this action and how to
 implement it",
 "time": "float, The time you decide to press the key",
 "key": "str, The key you decide to press, it must be one of the
 keys in your control config",
 "hold_for": "float, The duration you decide to press the key in
 seconds"
 }
]
]
]
}
```
# Control, Simulation, and Revision

- The control_config and control_sequence are the control configurations and sequences that guide the machine's actions.
- The simulation is the simulated motion trajectory of the machine, describing the motion trajectory (x, y, z) of some blocks in the machine.
- You should analyze the simulation and the control_config and control_sequence to revise the control_config and control_sequence to optimize the task.

```